



OPITZ CONSULTING

CONFIGURATION



INVENTORIES

- Contain **all hosts** that Ansible should know
- Hosts can be assigned to **groups**
- Can contain **variables**
(Best practice: outsource variables)
- Static Inventory Format
 - INI, YAML or JSON
- Inventory parameter can point to directory with many files

Default: /etc/ansible/hosts

```
localhost ansible_connection=local

[webservers]
node1.example.com ; Some comment
node2.example.com www_root=/var/http ; host var
node[3-10].example.com ;

[webservers:vars]
www_root=/var/www/ ; var for all webservers
```

USE OF GROUPS

- Primary addressing of hosts
 - How do I want to address my hosts?
- Meaningful groupings
 - Function (database, webserver, ...)
 - Location (dc1, dc2, frankfurt, berlin, ...)
 - Stage (develop, production)
 - Infrastructure type (vms, aws, ...)
- Multiple grouping possible
- Group hierarchies possible

[web]

```
web1 ansible_host=w0102.acme.com  
web2 ansible_host=w0103.acme.com
```

[db]

```
db1 ansible_host=w0104.acme.com  
db2 ansible_host=w0105.acme.com
```

[dev]

```
db1  
web1
```

[prod]

```
db2  
web2
```

SELECTION OF HOSTS

Example: `ansible [selection] -m ping`

Selection	Hosts
all	web1, web2, db1, db2
webservers1	web1, web2
webservers1:db	web1, web2, db1, db2
webservers1:&prod	web2
webservers1:db:&prod	web2, db2
webservers1[0]	web1
webservers1:! webservers2	web1

[webservers1]

web1 ansible_host=w0101.acme.com

web2 ansible_host=w0102.acme.com

[webservers2]

web2 ansible_host=w0102.acme.com

web3 ansible_host=w0103.acme.com

[db]

db1 ansible_host=w0104.acme.com

db2 ansible_host=w0105.acme.com

[dev]

web1

db1

[prod]

web2

db2

[atlanta] INI-Sample

host1
host2

[raleigh]

host2
host3

[southeast:children]

atlanta
raleigh

[southeast:vars]

some_server=foo.southeast.example.com
halon_system_timeout=30
self_destruct_countdown=60
escape_pods=2

[usa:children]

southeast
northeast
southwest
northwest

YAML-Sample

usa:

children:
southeast:

children:
atlanta:

hosts:
host1:
host2:

raleigh:

hosts:
host2:
host3:

vars:

some_server: foo.southeast.example.com
halon_system_timeout: 30
self_destruct_countdown: 60
escape_pods: 2

northeast:

northwest:

southwest:

DYNAMIC INVENTORIES

- Inventory is determined dynamically
- Useful for
 - Dynamic infrastructure (e.g. cloud provider, VMware)
 - Externally managed infrastructure (CMDB)
- Variant 1: Own inventory script (e.g. bash, python, ruby, ...)
 - Standardized interface
 - Understands parameters `--list` and `--host [host]`
 - Provides JSON-formatted inventory
- Variant 2: Inventory plug-in
 - Plugins deactivated by default (`enable_plugins` in Config)
 - Many available (`ansible-doc -t inventory -l`)

INVENTORY PLUGIN

nmap Uses nmap to find hosts to target

host_list Parses a 'host list' string

hcloud Ansible dynamic inventory plugin for the Hetzner Cloud

openstack OpenStack inventory source

vultr Vultr inventory source

aws_ec2 EC2 inventory source

cloudscale cloudscale.ch inventory source

virtualbox virtualbox inventory source

constructed Uses Jinja2 to construct vars and groups based on existing inventory

k8s Kubernetes (K8s) inventory source

azure_rm Azure Resource Manager inventory plugin

script Executes an inventory script that returns JSON

vmware_vm_inventory VMware Guest inventory source

openshift OpenShift inventory source

docker_machine Docker Machine inventory source

yaml Uses a specific YAML file as an inventory source

```
ansible-doc -t inventory -l
```

CONFIGURATION OF ANSIBLE

- What can be configured (among other things)
 - Path to inventory, path to additional roles
 - Number of forks (parallelism)
 - SSH options (e.g. sudo user, pipelining)
 - Dealing with "facts"
 - Etc.

CONFIGURATION OF ANSIBLE

- Sources for configuration (by ranking, will be aggregated)
 - Configuration settings
 - Configuration file
 - Environment variables
 - Command-line options
 - Playbook keywords
 - Variables

- https://docs.ansible.com/ansible/latest/reference_appendices/general_precedence.html#general-precedence-rules

PRECEDENCE EXAMPLE – BECOME (SUPERUSER)

- What did we learn:
 - Configuration settings
 - Command-line options
 - Playbook keywords
 - Variables

```
[privilege_escalation]
# (boolean) Toggles the use of privilege escalation, allowing you to 'become' another user after login.
become=True
```

```
ansible-playbook -b site.yml
```

```
---
- hosts: testkiste
  gather_facts: false
  become: false
  tasks:
    - name: test
      ansible.builtin apt:
        name: unzip
```

```
---
- hosts: testkiste
  gather_facts: false
  become: true
  become_user: sysadmin
  vars:
    ansible_become_user: admin
  tasks:
    - name: test
      ansible.builtin apt:
        name: unzip
```

- Result: become is set to false in the playbook run

ANSIBLE.CFG EXAMPLE VALUES

- Configuration file (Ansible picks the first it finds)

- ANSIBLE_CONFIG=/path/to/file
- ./ansible.cfg
- ~/.ansible.cfg
- /etc/ansible/ansible.cfg

- Default cfg ist generated with:

- ansible-config init --disabled > ansible.cfg

```
1  [defaults]
2  # (pathlist) Comma separated list of Ansible inventory sources
3  ;inventory=/etc/ansible/hosts
4
5  # (path) Option for connections using a certificate or key file to authenticate, rather than an age
6  ;private_key_file=
7
8  # (integer) Port to use in remote connections, when blank it will use the connection plugin default
9  ;remote_port=
10
11 # (string) Sets the login user for the target machines
12 # When blank it uses the connection plugin's default, normally the user currently executing Ansible
13 ;remote_user=
14
15 # (paths) Colon separated paths in which Ansible will search for Roles.
16 ;roles_path={{ ANSIBLE_HOME ~ "/roles:/usr/share/ansible/roles:/etc/ansible/roles" }}
17
18 [privilege_escalation]
19 # (boolean) Toggles the use of privilege escalation, allowing you to 'become' another user after lo
20 ;become=False
21
```